

Fall 2024 COMP 3511 Homework Assignment #1

Handout Date: September 15, 2024, Due Date: October 2, 2024

Name	
Student ID	
ITSC email	@connect.ust.hk

Please read the following instructions carefully before answering the questions:

- You must finish the homework assignment **individually**.
- This homework assignment contains **three** parts: (1) multiple choices, (2) short answer
3) programs with fork()
- **Homework Submission:** Please submit your homework to **Homework #1** on **Canvas**.
- TA responsible for HW1: Jingcan Chen, jchenhv@cse.ust.hk

1. [30 points] Multiple Choices

Write your answers in the boxes below:

MC1	MC2	MC3	MC4	MC5	MC6	MC7	MC8	MC9	MC10

1) Which of the following services is NOT necessary for an operating system?

- A) File System
- B) Error Detection
- C) GUI
- D) I/O Operation

2) Which of the following statements about *interrupt* is **not true**?

- A) Interrupts are used to handle asynchronous events, e.g., I/O ops and software errors
- B) Interrupts can have different importance specified by priority levels
- C) Upon each interrupt, a piece of OS codes is called and executed
- D) Interrupts cannot be generated by external hardware

3) Which of the following statements is **NOT true** on *process* and *thread*?

- A) At a certain moment, only one process is running on a CPU core
- B) One process with a single thread can run on different CPU cores at a time in parallel
- C) A process can consist of multiple thread(s), and they run within the same address space
- D) One process consists of only one address space

- 4) Which of the following statements is **not true** about the goal of *operating system*?
- (A) Execute user programs and make solving user problems easier
 - (B) Allow or deny user's access to hardware resources
 - (C) Make the computer system convenient to use
 - (D) Manage and use the computer hardware in an efficient manner
- 5) Which of the following statements is **true** about *system calls* and *dual mode operation*?
- A) Under dual mode operation, both user and the operating system can access all resources
 - B) The concept of dual modes can only have two modes
 - C) Some instructions are privileged in system calls
 - D) Context switch happens during system calls
- 6) Which of the following statements is **not true** about *loadable kernel module* approach?
- A) All components in the kernel can be dynamically added and removed during runtime
 - B) The kernel has a set of core components and can link in additional services via modules, either at boot time or during run time.
 - C) Module approach resembles a layered design in that each kernel section has a well-defined, protected interface, and it is flexible to call other modules
 - D) Loadable modules can add new services to the kernel without recompiling the whole kernel.
- 7) Which of the following statement is TRUE for *direct memory access* or DMA?
- A) DMA transfer data between an I/O device and memory directly without any assistance from CPU
 - B) DMA completes a data transfer without interrupting CPU
 - C) DMA frees up CPU from data movement between an I/O device and memory
 - D) All of the above
- 8) Which is true about *signal handling* in operating systems?
- A) Signals cannot be ignored and should terminate the program since they hurt system performance
 - B) Keyboard interrupt is a synchronous signal because it happens during the process execution, as the illegal memory access
 - C) Signals should be sent to every thread in a multi-threaded process
 - D) Each signal has a default handler function
- 9) When the degree of multiprogramming is too high, ____
- A) Long-term scheduler will remove some processes from the memory
 - B) Short-term scheduler assigns less CPU cores to a certain process
 - C) Mid-term scheduler swaps out certain processes to the disk
 - D) None of the above

TA responsible for HW1: Jingcan Chen (jchenhv@cse.ust.hk)

10) Which is **not true** about *interprocess communication (IPC)*?

- A) There are generally two models of IPC, i.e. shared memory and message passing
- B) Pipes can be used only in the parent-child processes or among multiple threads of a process
- C) Message sending need not block the processes
- D) Different processes communicating through sockets must have different port numbers

2. [30 points] Please answer the following questions in a few sentences

(1) (5 points) Please describe what the CPU needs to specify prior to DMA operations, and illustrate why this is better than programmed I/O when moving large chunks of data.

(2) (5 points) Please briefly explain the two essential properties (i.e., *spatial and temporal locality*) why caching works.

TA responsible for HW1: Jingcan Chen (jchenhv@cse.ust.hk)

(3) (5 points) What resources are shared and not shared in a multi-threaded process? What is the main benefit of using multiple threads instead of multiple processes?

(4) (5 points) What are the advantages of providing *system call APIs* to users in *dual-mode system*? Please explain from user view and system view.

(5) (5 points) What is *orphan* process? What is the problem with an *orphan* process? How does OS handle that?

(6) (5 points) What is *copy-on-write* ? What is the advantage of using *copy-on-write* in `fork()` implementation in Linux?

3. (40 points) Simple C programs on fork(). Suppose all printf() will be followed by fflush(stdout) even if it is not presented in the code. For all the C programs, you can assume that necessary header files are included

- 1) (5 points) Consider the following code segments:

```
int main()
{
    pid_t pid1;
    pid_t pid2;

    pid1 = fork();
    pid2 = fork();

    printf("pid1:%d, pid2:%d\n", pid1, pid2);
}
```

- (a) How many processes are there in total when this code finishes?
Give the answer.
- (b) If one process prints "pid1:51234, pid2: 51235", one process prints "pid1:0, pid2: 51236", write down other processes' pid and their outputs.

2) (15 points) Consider the following code segments:

```
int main() {
    int i = 0;
    int cnt = 10;
    for (; i < 3; i++) {
        pid_t pid = fork();
        if (pid == 0)
            printf("%d\n", cnt);
        else
            cnt += 10;
    }
    return 0;
}
```

- (a) Who will print the “cnt”, the child process or parent process? Why? How many times of “printf” will this code execute? Give the answer with explanation.

- (b) Theoretically, based on Linux fork() mechanism introduced in the lecture, how many memory copies of the variable “cnt” will there be? Explain your answer.

- (c) Directly running this piece of code may produce different outputs every time. Briefly explain why.

3) (10 points) Consider the following code segments:

```
#include <stdio.h>
#include <unistd.h>
int main() {
    for (int i = 0; i < 2; ++i) {
        if ((fork() && fork()) || !fork()) {
            printf("A");
            fflush(stdout);
        }
    }
}
```

- (a) Determine the total number of "A"s output by the given program, assuming it operates normally. Provide a brief explanation for your answer.

- 4) (10 points) Fill in the missing blanks using “printf” and “wait” functions, so that the following program will always display the following output:

Output: CDBA

Question:

```
int main() {
    if ( fork() ) {
        wait(0);
        if ( fork() ) {
            BLANK1;
            BLANK2;
            fflush(stdout);
        } else {
            printf("B");
            fflush(stdout);
        }
    } else {
        if ( !fork() ) {
            BLANK3;
            fflush(stdout);
        } else {
            BLANK4;
            printf("D");
            fflush(stdout);
        }
    }
    return 0;
}
```

BLANK1	
BLANK2	
BLANK3	
BLANK4	