

Fall 2024 COMP 3511 Homework Assignment #3

Handout Date: October 30 2024, Due Date: November 13 2024

Name	LI,Yuntong
Student ID	20944800
ITSC email	ylino@connect.ust.hk

Please read the following instructions carefully before answering the questions:

- You must finish the homework assignment **individually**.
- When you write your answers, please try to be precise and concise.
- **Homework Submission:** Please submit your homework to **Homework #3** on **Canvas**.
- TA responsible for HW3: Xingxing TANG (xtangav@connect.ust.hk)

1. (30 points) Multiple Choices

Write your answers in the boxes below:

MC1	MC2	MC3	MC4	MC5	MC6	MC7	MC8	MC9	MC10
D	C	C	C	D	D	D	D	A	A

(1) Which of the following statements is TRUE about *Race Condition*?

- A) Race Condition can cause potential data inconsistency
- B) The outcome of the executions depends on the particular order
- C) Even race condition exists, the outcome of the executions can be correct
- D) All of the above

(2) Which of the following statements is TRUE about *Mutex Lock*?

- A) Calling acquire() of a mutex lock can be interrupted
- B) Calling release() of a mutex lock can be interrupted
- C) Mutex lock is often implemented as a spinlock
- D) All of the above

(3) Which of the following statements is FALSE about *Semaphore*?

- A) Improper usage of semaphore can cause deadlock
- B) Semaphore can be used as a mutex lock
- C) Semaphore values cannot become negative
- D) Semaphore has busy waiting and non-busy waiting versions

(4) Which of the following statements is TRUE about *Reader-Writer Problems*?

- A) Multiple readers cannot read the shared data simultaneously

TA responsible for HW3: Xingxing TANG (xtangav@connect.ust.hk)

- B) We should always prioritize readers in reader-writer problem
- C) OS usually provides a reader-writer lock to mitigate the starvation problem
- D) None of the above

(5) Which of the following statements is TRUE about *Deadlock*?

- A) Deadlock prevention and avoidance guarantee that deadlock will never happen
- B) Circular wait always leads to deadlock
- C) The number of resource instances has no impact on whether the deadlock will happen
- D) None of the above

(6) Which of the following statement is TRUE about the *Bankers' Algorithm*?

- A) Bankers' algorithm does not need to know the maximum resource usage for each process
- B) The Safety Algorithm is only invoked once at the beginning of the algorithm
- C) After resource is allocated using Resource-Request Algorithm, we never restore the state of resource allocation
- D) None of the above

(7) Which of the following statement is FALSE about the *Deadlock Detection Algorithm*?

- A) The deadlock detection algorithm is identical to the safety algorithm
- B) Using deadlock detection algorithm for each resource allocation request causes large overhead
- C) The deadlock detection algorithm can tell which process cause the deadlock
- D) All of the above

(8) Consider a segment table of one process:

Segment #	Segment length	Starting address	Permission	Status
0	100	6000	Read-only	In memory
1	150	5500	Read/Write	In memory
2	350	4000	Read/Write	In memory

When accessing the logical address at <segment # = 2, offset = 400>, the results after address translation is:

- A) No such segment
- B) Return address 4400
- C) Invalid permission
- D) Address out of range

(9) In a system with 32-bit address, given the logical address 0x0000F1BC (in hexadecimal) with a page size of 256 bytes, what is the page offset?

- A) 0xBC
- B) 0xF1

TA responsible for HW3: Xingxing TANG (xtangav@connect.ust.hk)

C) 0xC

D) 0xF100

(10) Which of the following statement is FALSE when comparing Segmentation and Paging Schemes?

A) The number of entries in a page table is usually smaller than the number of entries in a segmentation table

B) Paging scheme results in better memory utilization than segmentation scheme

C) Sharing a segment is easier than sharing a page, as each segment represents a logical entity in a program

D) Paging scheme does not suffer from external fragmentation

2. (20 points) Process Synchronization

This problem involves multiple readers and one writer accessing a shared resource.

When the writer wants to write, readers must wait, ensuring the writer gets priority.

We will use a new API `pthread_cond_broadcast(&cond)` to wake up all waiting threads on the conditional variable `cond`.

Please fill in 10 blanks to correctly implement synchronization using mutex locks and condition variables.

```
#include <pthread.h>
```

```
#include <stdio.h>
```

```
#include <unistd.h>
```

```
int read_count = 0;    // Number of active readers
```

```
int writer_waiting = 0; // Is the writer waiting?
```

```
pthread_mutex_t read_count_mutex; // Protects the read_count variable
```

```
pthread_mutex_t writer_mutex;    // Controls writer's priority and access
```

```
pthread_cond_t reader_cond;      // Condition variable for readers
```

```
pthread_cond_t writer_cond;      // Condition variable for the writer
```

```
void* reader(void* arg) {
```

```
    // Lock to modify read_count
```

```
    pthread_mutex_lock(&read_count_mutex); // (1) Lock read_count_mutex to access read_count
```

```
    // Wait if the writer is waiting
```

```
    while (writer_waiting > 0) { // (2) Check if the writer is waiting
```

```
        pthread_cond_wait(&reader_cond, &read_count_mutex); // (3) Wait for the writer
```

```
    }
```

```
    read_count++; // Increment reader count
```

TA responsible for HW3: Xingxing TANG (xtangav@connect.ust.hk)

```
pthread_mutex_unlock(&read_count_mutex); // Unlock read_count_mutex

// Reading section
printf("Reader %d is reading.\n", id);

// Lock to decrement read_count
pthread_mutex_lock(&read_count_mutex); // Lock read_count_mutex
read_count--; // Decrement reader count

// If no readers are active, signal the writer
if (read_count == 0) { // (4) Check if no readers are active
    pthread_cond_signal(&writer_cond); // (5) Signal the writer
}
pthread_mutex_unlock(&read_count_mutex); // Unlock read_count_mutex

return NULL;
}

void* writer(void* arg) {
    // Lock to manage writer state and access
    pthread_mutex_lock(&writer_mutex); // (6) Lock writer_mutex to access writer state
    writer_waiting = 1; // Mark the writer as waiting

    // Wait until all readers are done
    pthread_mutex_lock(&read_count_mutex); // Lock read_count_mutex
    while (read_count > 0) { // (7) Check if any readers are active
        pthread_cond_wait(&writer_cond, &read_count_mutex); // (8) Wait for readers to finish
    }
    pthread_mutex_unlock(&read_count_mutex); // Unlock read_count_mutex

    // Writing section
    printf("Writer %d is writing.\n", id);
    writer_waiting = 0; // Mark the writer as not waiting

    // Signal waiting readers to proceed
    pthread_cond_broadcast(&reader_cond); // (9) Broadcast to readers
    pthread_mutex_unlock(&writer_mutex); // (10) Unlock writer_mutex

    return NULL;
}
```

}

3. (30 points) Deadlocks

Consider the following snapshot of a system:

	Allocation					Max					Available			
	A	B	C	D		A	B	C	D		A	B	C	D
P0	2	0	0	1		4	2	3	3		2	3	2	1
P1	4	1	2	1		6	1	3	2					
P2	1	1	0	3		2	4	1	6					
P3	1	2	1	2		1	3	2	4					
P4	1	4	5	1		3	4	6	1					

(1) (5 points) What is the content of the Need matrix?

```

      A B C D
P0  2 2 3 2
P1  2 0 1 1
P2  1 3 1 3
P3  0 1 1 2
P4  2 0 1 0

```

(2) (10 points) Is the system in a safe state? Why?

Yes, the system is in a safe state.

Work = Available resources = [2, 3, 2, 1]

• P0: Need = [2, 2, 3, 2] > Work

• P1: Need = [2, 0, 1, 1] ≤ Work

Work = [2, 3, 2, 1] + [4, 1, 2, 1] = [6, 4, 4, 2]

• P2: Need = [1, 3, 1, 3] > Work

• P3: Need = [0, 1, 1, 2] ≤ Work

Work = [6, 4, 4, 2] + [1, 2, 1, 2] = [7, 6, 5, 4]

• P4: Need = [2, 0, 1, 0] ≤ Work

Work = [7, 6, 5, 4] + [1, 4, 5, 1] = [8, 10, 10, 5]

• P2: Need = [1, 3, 1, 3] ≤ Work

Work = [8, 10, 10, 5] + [1, 1, 0, 3] = [9, 11, 10, 8]

• P0: Need = [2, 2, 3, 2] ≤ Work

Work = [9, 11, 10, 8] + [2, 0, 0, 1] = [11, 11, 10, 9]

Since we have found a sequence in which all processes can finish (P1 → P3 → P4 → P2 → P0), the system is in a safe state.

(3) (5 points) If a request from process P1 arrives for (1, 1, 0, 2), can the request be granted? Why?

No, because $(1, 1, 0, 2) > \text{Need}[1] = (2, 0, 1, 1)$ at resource B
So this request is not allowed, and process P1 should be blocked.

(4) (10 points) If a request from process P0 arrives for (1, 1, 2, 1), can the request be granted? Why?

The request cannot be granted.

The request $[1, 1, 2, 1] \leq \text{Need} [2, 2, 3, 2]$, and $[1, 1, 2, 1] \leq \text{Available} [2, 3, 2, 1]$

Next,

- the Updated Allocation Matrix: $[[3, 1, 2, 2] [4, 1, 2, 1] [1, 1, 0, 3] [1, 2, 1, 2] [1, 4, 5, 1]]$
- the Updated Need Matrix: $[[1, 1, 1, 1] [2, 0, 1, 1] [1, 3, 1, 3] [0, 1, 1, 2] [2, 0, 1, 0]]$

Since Available Resources $[1, 2, 0, 0]$ is smaller than any Need from each process, the request from process P0 for (1, 1, 2, 1) cannot be granted.

TA responsible for HW3: Xingxing TANG (xtangav@connect.ust.hk)

4. (20 points) Memory Management

Consider the segment table shown in Table A. Translate each of the virtual addresses in Table B into physical addresses. Indicate errors (out of range, no such segment) if an address cannot be translated.

Table A

Segment number	Starting address	Segment length
0	260	90
1	1466	160
2	2656	130
3	146	50
4	2064	370

Table B

Segment number	Offset
0	10
1	180
2	100
3	80
4	50
5	32

Answer:

Segment number	Offset	Physical address
0	10	270
1	180	Out of range
2	100	2756
3	80	Out of range
4	50	2114
5	32	No such segment

5. (12 points) Paging

Consider a virtual memory system providing 256 pages for each user program; the size of each page is 4KB. The size of main memory is 256KB. Consider one user program occupied 4 pages, and the page table of this program is shown as below:

Logical page number	Physical block number
0	9
1	6
2	7
3	3

Assume there are two requests on logical address (in hexadecimal number) 010AF, 100AF.

- (1) (4 points) Please describe how many bits the virtual page number and the page offset contain, respectively.

Virtual page number: 8 bits (since $256 \text{ pages} = 2^8$)

Page offset: 12 bits (since $4\text{KB} = 2^{12} \text{ bytes}$)

- (2) (8 points) Please illustrate how the virtual memory system will deal with these requests. (Please indicate if there will be a page fault. If there is no page fault, please give the corresponding physical address.)

(010AF):

Binary: 0000 0001 0000 1010 1111

Virtual Page Number: 00000001

1 exists in Page Table, maps to Physical Block Number 6.

Page Offset: 0000 1010 1111 (0x0AF)

The physical address will be 0x060AF.

(100AF): Page fault occurs, as page number 16 is not mapped.

TA responsible for HW3: Xingxing TANG (xtangav@connect.ust.hk)