


Chapter 17: Protection





Chapter 17: Protection

- Goals of Protection
- Principles of Protection
- Protection Rings
- Domain of Protection
- Access Matrix
- Implementation of Access Matrix





Objectives

- ❑ Discuss the goals and principles of protection in modern computer systems
- ❑ Explain how **protection domains** combined with **an access matrix** are used to specify the resources that a process may access
- ❑ Examine capability-based protection system





Goals of Protection

- In a protection model, computer system consists of a collection of **objects**, hardware or software
 - Hardware objects: CPU, memory segments, printers, disks, and tape
 - Software objects: files, programs, and semaphores
- Each object has a unique name and can be accessed through a **well-defined set of operations**
- **Protection problem** is to ensure that each object is accessed correctly and only by those processes allowed to do so
- **Mechanisms** are distinct from **policies**, in which mechanisms determine how something will be done, and policies decide what will be done.
 - The separation is important for **flexibility**, as policies are likely to change from place to place or from time to time.
 - The separation ensures that not every change in policy would require a change in the underlying mechanism.





Principles of Protection


- The guiding principle – **principle of least privilege**
 - Programs, users and systems should be given just enough privileges to perform their tasks - mitigate the attack
 - In file permissions, this principle dictates that a user have read access but not write or execute access to a file. **The principle of least privilege** would require that the OS provides a mechanism to only allow read access but not write or execute access
- Properly set **permissions** (i.e., the access rights to an object) can limit damage if entity has a bug or gets abused





Protection Rings

- ❑ **User mode** and **kernel mode** – **privilege separation**
- ❑ Hardware support required to support the notion of separate execution
- ❑ Let D_i and D_j be any two domain rings
- ❑ If $j < i \Rightarrow D_i \subseteq D_j$
- ❑ The innermost ring, ring 0, provides the full set of privileges





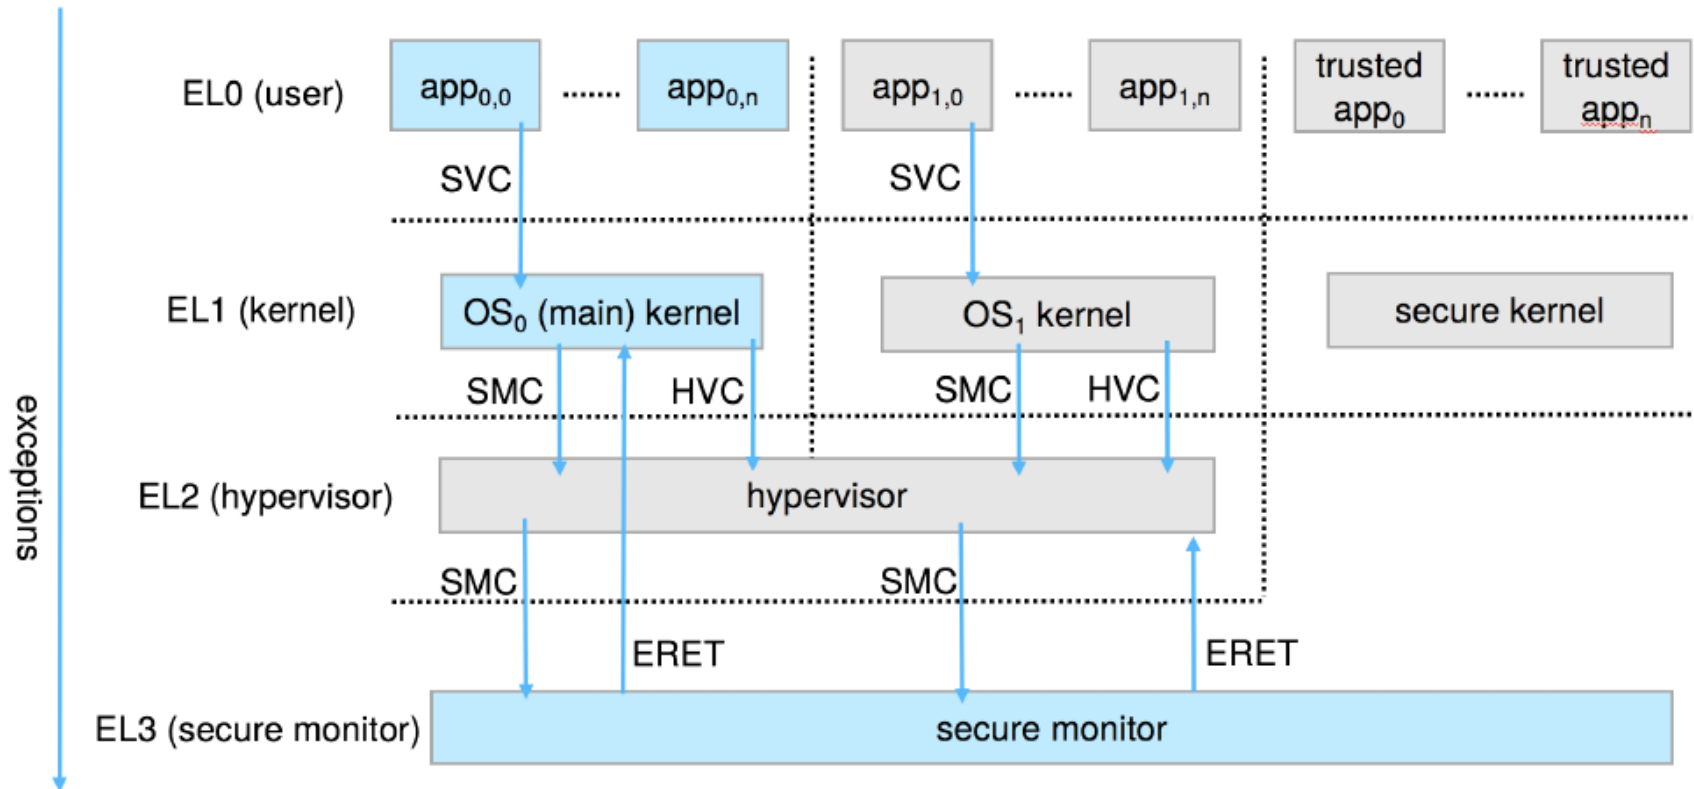
Protection Rings (Cont.)

- ❑ Components ordered by amount of privilege and protected from each other
 - ❑ For example, the kernel is in one ring and user applications in another
 - ❑ This privilege separation requires hardware support
 - ❑ “Gates” used to transfer between rings, for example the **syscall** Intel instruction, also **traps** and **interrupts**
- ❑ **Hypervisors** (Intel) is introduced (another ring) - virtual machine managers, which create and run virtual machines, and have more capabilities than the kernels of the guest operating systems
- ❑ ARM processors added **TrustZone** or **TZ** ring to protect crypto functions with access (more privileged than kernel)
 - ❑ This most privileged execution environment has exclusive access to hardware-backed cryptographic features, such as the NFC Secure Element and an on-chip cryptographic key, that make handling passwords and sensitive information more secure.





ARM CPU Architecture





Domain of Protection

- ❑ Protection rings separate functions into different domains and order them hierarchically
- ❑ **Domain** can be considered as a generalization of rings without a [hierarchy](#)
- ❑ A computer system can be treated as processes and objects
 - ❑ **Hardware objects** (such as CPU, memory, disk) and **software objects** (such as files, programs, semaphores)
- ❑ Process for example should only have access to objects it currently requires to complete its task – the [need-to-know](#) principle (policy)
- ❑ Implementation can be via process operating in a [protection domain](#)
 - ❑ Protection domain specifies the set of resources a process may access
 - ❑ Each domain specifies set of objects and types of operations may be invoked on each object





Domain of Protection (Cont.)

- The ability to execute an operation on an object is an **access right**
- A **domain** is a collection of access rights, each of which is an ordered pair **<object-name, rights-set>**
 - An example: if domain D has the access right <file F, {read,write}>, then a process executing in domain D can both read and write file F. It cannot, however, perform any other operation on that object.
- Domains may share access rights
- Associations between processes and domains can be **static** if the set of resources available to the process is fixed throughout the process's lifetime, or can be **dynamic**
- If the association is dynamic, a mechanism is available to allow **domain switching**, enabling the process to switch from one domain to another during different stage of execution





Domain of Protection (Cont.)

Domain can be realized in a variety of ways:

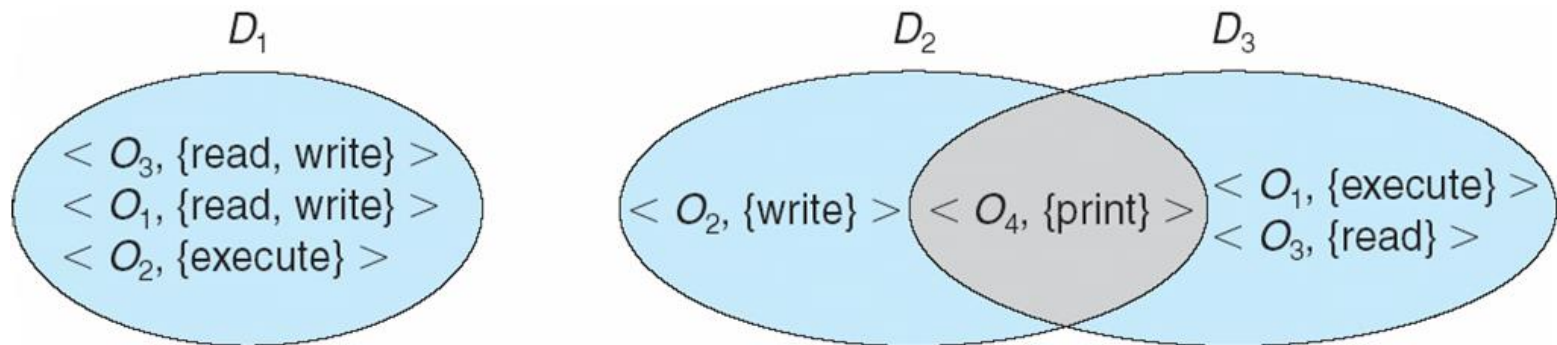
- ❑ Each **user** may be a domain - the set of objects that can be accessed depends on the identity of the user. Domain switching occurs when the user is changed
- ❑ Each **process** may be a domain - the set of objects that can be accessed depends on the identity of the process. Domain switching occurs when one process sends a message to another process and then waits for a response.
- ❑ Each **procedure** may be a domain - the set of objects that can be accessed corresponds to the local variables defined within the procedure. Domain switching occurs when a procedure call is made





Domain Structure

- **Access-right** = $\langle \text{object-name}, \text{rights-set} \rangle$
where rights-set is a subset of all valid operations that can be performed on the object
- **Domain** = set of access-rights
- The access right $\langle O_4, \{\text{print}\} \rangle$ shared by domains D_2 and D_3 , thus, a process executing in either of these two domains can print object O_4 .





Access Matrix

- View protection as a matrix (**access matrix**)
- **Rows** represent domains, and **columns** represent objects
- **Access(i,j)** consists of a set of access rights - the set of operations that a process executing in Domain_i can invoke on Object_j

object domain	F_1	F_2	F_3	printer
D_1	read		read	
D_2				print
D_3		read	execute	
D_4	read write		read write	





Use of Access Matrix

- The **access matrix** scheme provides the mechanism for specifying a variety of policies - mechanism and policy separation
- The **mechanism** consists of implementing the access matrix and ensuring that the semantic properties hold.
 - To ensure that a process executing in domain D_i can access only those objects specified in row i .
- The **policy** decisions specify which rights should be included in the (i,j) th entry, and determine the domain in which each process executes
- If a process in Domain D_i tries to do “op” on object O_j , then “op” must be in the access matrix
- User who creates an object can define access column for that object
 - When a user creates a new object O_j , the column O_j is added to the access matrix with the appropriate initialization entries, as dictated by the creator. The user may decide to enter some rights in some entries in column j and other rights in other entries, as needed.





Use of Access Matrix (Cont.)

This can be expanded to dynamic protection

- Operations to add, delete access rights
- Special access rights:
 - **owner** of O_i - can add and remove any right in any entry in column
 - **copy** op from O_i to O_j (denoted by “*”) - only within the column (that is, for the object)
 - **control** – D_i can modify D_j access rights – modify domain objects (a row)
 - **transfer** – switch from domain D_i to D_j
- **Copy** and **Owner** applicable to an object - change the entries in a column
- **Control** applicable to domain object - change the entries in a row
- New objects and new domains can be created dynamically and included in the access-matrix model
- In a **dynamic** protection system, we may sometimes need to revoke access rights to objects shared by different users – **revocation** of access right





Access Matrix of Figure A with Domains as Objects

domain \ object	F_1	F_2	F_3	laser printer	D_1	D_2	D_3	D_4
D_1	read		read			switch		
D_2				print			switch	switch
D_3		read	execute					
D_4	read write		read write		switch			





Access Matrix with Copy Rights

object domain	F_1	F_2	F_3
D_1	execute		write*
D_2	execute	read*	execute
D_3	execute		

(a)

object domain	F_1	F_2	F_3
D_1	execute		write*
D_2	execute	read*	execute
D_3	execute	read	

(b)





Access Matrix With Owner Rights

object domain	F_1	F_2	F_3
D_1	owner execute		write
D_2		read* owner	read* owner write
D_3	execute		

(a)

object domain	F_1	F_2	F_3
D_1	owner execute		write
D_2		owner read* write*	read* owner write
D_3		write	write

(b)





Modified Access Matrix of Figure B

object domain	F_1	F_2	F_3	laser printer	D_1	D_2	D_3	D_4
D_1	read		read			switch		
D_2				print			switch	switch control
D_3		read	execute					
D_4	write		write		switch			





Implementation of Access Matrix

- In general, the access matrix is sparse; that is, most of the entries will be empty
- **Option 1 – Global Table**
 - Store ordered triples $\langle \text{domain}, \text{object}, \text{rights-set} \rangle$ in table
 - A requested operation M on object O_j within domain $D_i \rightarrow$ search table for $\langle D_i, O_j, R_k \rangle$ with $M \in R_k$
 - But the table could be large \rightarrow might not fit in main memory, requires additional I/O – virtual memory techniques are often used
 - Difficult to group objects - For example, if everyone can read a particular object, this object must have a separate entry in every domain.





Implementation of Access Matrix (Cont.)

- Each column = **Access-control list** for one object
Defines who can perform what operation

Domain 1 = Read, Write

Domain 2 = Read

Domain 3 = Read

- Each row = **Capability List** (like a key)
For each domain, what operations allowed on what objects

Object F1 – Read

Object F4 – Read, Write, Execute

Object F5 – Read, Write, Delete, Copy





Implementation of Access Matrix (Cont.)

□ Option 2 – Access lists for objects

- Each column implemented as an access list for one object
- Resulting per-object list consists of ordered pairs $\langle \text{domain}, \text{rights-set} \rangle$ defining all domains with non-empty set of access rights for the object
- Obviously, the empty entries can be discarded.
- This can be easily extended to define **default** set of access rights -> If $M \in \text{default set}$, also allow access (for all domains)





Implementation of Access Matrix (Cont.)

□ Option 3 – Capability list for domains

- Instead of object-based, list is domain-based
- A **capability list** for domain is a list of objects together with operations allowed on them
- An object represented by its name or address, called a **capability**
- To execute operation M on object O_j , a process requests operation M , specifying the capability (or pointer) for object O_j as a parameter
- Possession of capability means access is allowed
- Capability list associated with a domain, but never directly accessible by a process executing in that domain
 - Rather, the capability list itself is a **protected object**, maintained by OS and accessed by users only indirectly
 - This avoids the possibility of capability list modification by users
 - If all capabilities are secure, the object they protect is also secure against unauthorized access





Implementation of Access Matrix (Cont.)

❑ Option 4 – Lock-key

- ❑ Compromise between access lists and capability lists
- ❑ Each object has list of unique bit patterns, called **locks**
- ❑ Each domain as list of unique bit patterns called **keys**
- ❑ Process in a domain can only access object if domain has key that matches one of the locks of the object
- ❑ As with capability lists, the list of keys for a domain must be managed by the operating system on behalf of the domain.
- ❑ Users are not allowed to examine or modify the list of keys (or locks) directly.





Comparison of Implementations

Choosing a technique for implementing an access matrix involves various trade-offs.

- ❑ **Global table** is simple, but large, lack of grouping of objects or domains
- ❑ **Access lists** correspond directly to the needs of users
 - ❑ An access list on an object is specified when a user creates the object
 - ❑ Determining set of access rights for each domain is difficult - every access to the object must be checked, requiring a search of the access list.
- ❑ **Capability lists** useful for localizing information for a given process
 - ❑ But revocation capabilities can be inefficient
- ❑ **Lock-key** can be effective and flexible depending on the length of the keys
 - ❑ Keys can be passed freely from domain to domain, easy revocation
- ❑ Most systems use combination of access lists and capabilities



End of Chapter 17

